

```
In [1]: import pandas as pd

In [2]: df = pd.read_csv("sales_data_sample.csv", encoding='latin1')

In [3]: df

Out[3]:
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	...
0	10107	30	95.70	2	2871.00	2/24/2003 0:00	Shipped	1	2	2003	...
1	10121	34	81.35	5	2765.90	5/7/2003 0:00	Shipped	2	5	2003	...
2	10134	41	94.74	2	3884.34	7/1/2003 0:00	Shipped	3	7	2003	...
3	10145	45	83.26	6	3746.70	8/25/2003 0:00	Shipped	3	8	2003	...
4	10159	49	100.00	14	5205.27	10/10/2003 0:00	Shipped	4	10	2003	...
...
2818	10350	20	100.00	15	2244.40	12/2/2004 0:00	Shipped	4	12	2004	...
2819	10373	29	100.00	1	3978.51	1/31/2005 0:00	Shipped	1	1	2005	...
2820	10386	43	100.00	4	5417.57	3/1/2005 0:00	Resolved	1	3	2005	...
2821	10397	34	62.24	1	2116.16	3/28/2005 0:00	Shipped	1	3	2005	...
2822	10414	47	65.52	9	3079.44	5/6/2005 0:00	On Hold	2	5	2005	...

2823 rows × 25 columns

```
In [4]: df.describe()
```

Out[4]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	MSRP
count	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000
mean	10258.725115	35.092809	83.658544	6.466171	3553.889072	2.717676	7.092455	2003.81509	100.715551
std	92.085478	9.741443	20.174277	4.225841	1841.865106	1.203878	3.656633	0.69967	40.187912
min	10100.000000	6.000000	26.880000	1.000000	482.130000	1.000000	1.000000	2003.00000	33.000000
25%	10180.000000	27.000000	68.860000	3.000000	2203.430000	2.000000	4.000000	2003.00000	68.000000
50%	10262.000000	35.000000	95.700000	6.000000	3184.800000	3.000000	8.000000	2004.00000	99.000000
75%	10333.500000	43.000000	100.000000	9.000000	4508.000000	4.000000	11.000000	2004.00000	124.000000
max	10425.000000	97.000000	100.000000	18.000000	14082.800000	4.000000	12.000000	2005.00000	214.000000

```
In [5]: df.isnull().sum()
```

Out[5]:

ORDERNUMBER	0
QUANTITYORDERED	0
PRICEEACH	0
ORDERLINENUMBER	0
SALES	0
ORDERDATE	0
STATUS	0
QTR_ID	0
MONTH_ID	0
YEAR_ID	0
PRODUCTLINE	0
MSRP	0
PRODUCTCODE	0
CUSTOMERNAME	0
PHONE	0
ADDRESSLINE1	0
ADDRESSLINE2	2521
CITY	0
STATE	1486
POSTALCODE	76
COUNTRY	0
TERRITORY	1074
CONTACTLASTNAME	0
CONTACTFIRSTNAME	0
DEALSIZE	0

dtype: int64

In [6]:

df.drop(columns=['ADDRESSLINE1','ADDRESSLINE2','STATUS','POSTALCODE','STATE','TERRITORY','CITY','ORDERNUMBER','CUSTOMER

In [7]:

df

Out[7]:

	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	PRODUCTLINE	MSRP	PRODUCTCODE	CO
0	30	95.70	2	2871.00	1	2	2003	Motorcycles	95	S10_1678	
1	34	81.35	5	2765.90	2	5	2003	Motorcycles	95	S10_1678	
2	41	94.74	2	3884.34	3	7	2003	Motorcycles	95	S10_1678	
3	45	83.26	6	3746.70	3	8	2003	Motorcycles	95	S10_1678	
4	49	100.00	14	5205.27	4	10	2003	Motorcycles	95	S10_1678	
...
2818	20	100.00	15	2244.40	4	12	2004	Ships	54	S72_3212	
2819	29	100.00	1	3978.51	1	1	2005	Ships	54	S72_3212	
2820	43	100.00	4	5417.57	1	3	2005	Ships	54	S72_3212	
2821	34	62.24	1	2116.16	1	3	2005	Ships	54	S72_3212	
2822	47	65.52	9	3079.44	2	5	2005	Ships	54	S72_3212	

2823 rows × 12 columns

In [8]:

DEALSIZE = pd.get_dummies(df['DEALSIZE'])

In [9]:

DEALSIZE

Out[9]:

	Large	Medium	Small
0	False	False	True
1	False	False	True
2	False	True	False
3	False	True	False
4	False	True	False
...
2818	False	False	True
2819	False	True	False
2820	False	True	False
2821	False	False	True
2822	False	True	False

2823 rows × 3 columns

In [10]:

PRODUCTLINE = pd.get_dummies(df['PRODUCTLINE'])
COUNTRY = pd.get_dummies(df['COUNTRY'])

In [11]:

df = pd.concat([df,DEALSIZE,PRODUCTLINE,COUNTRY], axis=1)

In [12]:

df

Out[12]:

	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	PRODUCTLINE	MSRP	PRODUCTCODE	...
0	30	95.70	2	2871.00	1	2	2003	Motorcycles	95	S10_1678	...
1	34	81.35	5	2765.90	2	5	2003	Motorcycles	95	S10_1678	...
2	41	94.74	2	3884.34	3	7	2003	Motorcycles	95	S10_1678	...
3	45	83.26	6	3746.70	3	8	2003	Motorcycles	95	S10_1678	...
4	49	100.00	14	5205.27	4	10	2003	Motorcycles	95	S10_1678	...
...
2818	20	100.00	15	2244.40	4	12	2004	Ships	54	S72_3212	...
2819	29	100.00	1	3978.51	1	1	2005	Ships	54	S72_3212	...
2820	43	100.00	4	5417.57	1	3	2005	Ships	54	S72_3212	...
2821	34	62.24	1	2116.16	1	3	2005	Ships	54	S72_3212	...
2822	47	65.52	9	3079.44	2	5	2005	Ships	54	S72_3212	...

2823 rows × 41 columns

```
In [13]: df.drop(columns=['DEALSIZE', 'PRODUCTLINE', 'COUNTRY'],axis=1, inplace=True)
```

```
In [14]: df
```

Out[14]:

	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	MSRP	PRODUCTCODE	Large	...	Italy	J&
0	30	95.70	2	2871.00	1	2	2003	95	S10_1678	False	...	False	F
1	34	81.35	5	2765.90	2	5	2003	95	S10_1678	False	...	False	F
2	41	94.74	2	3884.34	3	7	2003	95	S10_1678	False	...	False	F
3	45	83.26	6	3746.70	3	8	2003	95	S10_1678	False	...	False	F
4	49	100.00	14	5205.27	4	10	2003	95	S10_1678	False	...	False	F
...
2818	20	100.00	15	2244.40	4	12	2004	54	S72_3212	False	...	False	F
2819	29	100.00	1	3978.51	1	1	2005	54	S72_3212	False	...	False	F
2820	43	100.00	4	5417.57	1	3	2005	54	S72_3212	False	...	False	F
2821	34	62.24	1	2116.16	1	3	2005	54	S72_3212	False	...	False	F
2822	47	65.52	9	3079.44	2	5	2005	54	S72_3212	False	...	False	F

2823 rows × 38 columns



```
In [15]: df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes
```

```
In [16]: df
```

Out[16]:

	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	MSRP	PRODUCTCODE	Large	...	Italy	J&
0	30	95.70	2	2871.00	1	2	2003	95	0	False	...	False	F
1	34	81.35	5	2765.90	2	5	2003	95	0	False	...	False	F
2	41	94.74	2	3884.34	3	7	2003	95	0	False	...	False	F
3	45	83.26	6	3746.70	3	8	2003	95	0	False	...	False	F
4	49	100.00	14	5205.27	4	10	2003	95	0	False	...	False	F
...
2818	20	100.00	15	2244.40	4	12	2004	54	108	False	...	False	F
2819	29	100.00	1	3978.51	1	1	2005	54	108	False	...	False	F
2820	43	100.00	4	5417.57	1	3	2005	54	108	False	...	False	F
2821	34	62.24	1	2116.16	1	3	2005	54	108	False	...	False	F
2822	47	65.52	9	3079.44	2	5	2005	54	108	False	...	False	F

2823 rows × 38 columns



```
In [17]: df.dtypes
```

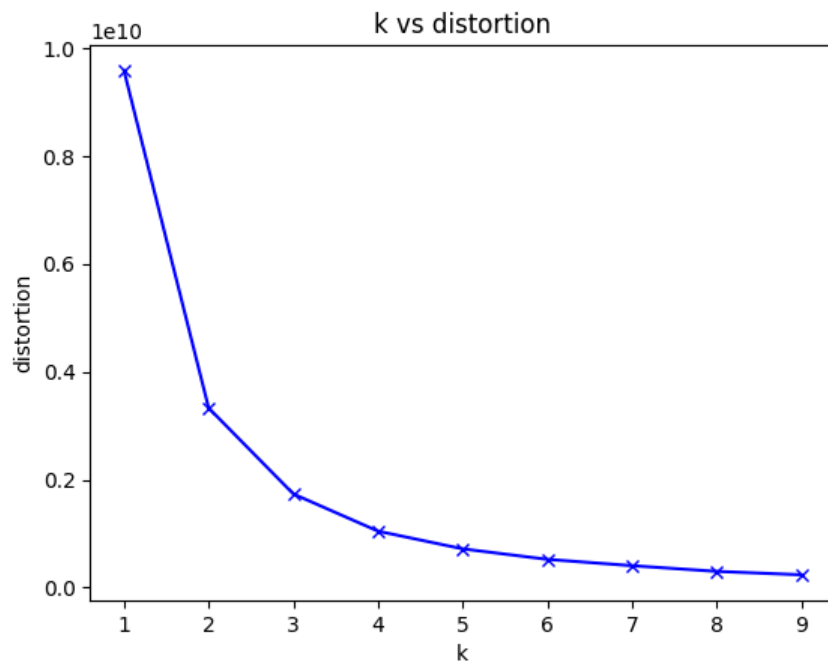
```
Out[17]: QUANTITYORDERED      int64
PRICEEACH                    float64
ORDERLINENUMBER              int64
SALES                        float64
QTR_ID                       int64
MONTH_ID                     int64
YEAR_ID                      int64
MSRP                         int64
PRODUCTCODE                  int8
Large                        bool
Medium                       bool
Small                        bool
Classic Cars                 bool
Motorcycles                  bool
Planes                       bool
Ships                        bool
Trains                       bool
Trucks and Buses             bool
Vintage Cars                 bool
Australia                    bool
Austria                      bool
Belgium                      bool
Canada                       bool
Denmark                      bool
Finland                      bool
France                       bool
Germany                      bool
Ireland                      bool
Italy                        bool
Japan                        bool
Norway                       bool
Philippines                  bool
Singapore                    bool
Spain                        bool
Sweden                       bool
Switzerland                  bool
UK                           bool
USA                          bool
dtype: object
```

```
In [18]: from sklearn.cluster import KMeans
```

```
In [19]: distortion = []
k = range(1,10)
for n in k:
    km = KMeans(n_clusters=n)
    km.fit(df)
    distortion.append(km.inertia_)
```

```
In [20]: import matplotlib.pyplot as plt
```

```
In [21]: plt.plot(k,distortion,'bx-')
plt.xlabel('k')
plt.ylabel('distortion')
plt.title('k vs distortion')
plt.show()
```



```
In [22]: x_train = df.values
model = KMeans(n_clusters=4, random_state=2)
model.fit(x_train)
```

```
Out[22]: KMeans(n_clusters=4, random_state=2)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [23]: pred = model.predict(x_train)
```

```
In [24]: print(pred)
```

```
[3 3 3 ... 2 0 3]
```

```
In [25]: import numpy as np
```

```
In [26]: unique,count = np.unique(pred, return_counts=True)
```

```
In [27]: print(unique)
```

```
[0 1 2 3]
```

```
In [28]: print(count)
```

```
[1041  199  562 1021]
```

```
In [29]: pred1 = pd.DataFrame(pred)
```

```
In [30]: pred1
```

Out[30]:

	0
0	3
1	3
2	3
3	3
4	2
...	...
2818	0
2819	3
2820	2
2821	0
2822	3

2823 rows × 1 columns

```
In [31]: df = pd.concat([df,pred1],axis=1)
df
```

Out[31]:

	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	MSRP	PRODUCTCODE	Large	...	Japan	...
0	30	95.70	2	2871.00	1	2	2003	95	0	False	...	False	...
1	34	81.35	5	2765.90	2	5	2003	95	0	False	...	False	...
2	41	94.74	2	3884.34	3	7	2003	95	0	False	...	False	...
3	45	83.26	6	3746.70	3	8	2003	95	0	False	...	False	...
4	49	100.00	14	5205.27	4	10	2003	95	0	False	...	False	...
...
2818	20	100.00	15	2244.40	4	12	2004	54	108	False	...	False	...
2819	29	100.00	1	3978.51	1	1	2005	54	108	False	...	False	...
2820	43	100.00	4	5417.57	1	3	2005	54	108	False	...	False	...
2821	34	62.24	1	2116.16	1	3	2005	54	108	False	...	False	...
2822	47	65.52	9	3079.44	2	5	2005	54	108	False	...	False	...

2823 rows × 39 columns



```
In [35]: df_temp = df.drop(columns=[0],axis=1)
```

```
In [39]: #cluster_centers is used to find the centroid values of the clusters
model.cluster_centers_
```

```
Out[39]: array([[ 2.98893167e+01,  6.54380366e+01,  6.61886429e+00,
  1.88419459e+03,  2.72473532e+00,  7.13282002e+00,
  2.00381521e+03,  7.33378248e+01,  6.32569779e+01,
  2.08166817e-17,  5.55111512e-16,  1.00000000e+00,
  2.59865255e-01,  1.22232916e-01,  1.17420597e-01,
  8.95091434e-02,  4.42733397e-02,  8.75842156e-02,
  2.79114533e-01,  7.21847931e-02,  1.73243503e-02,
  1.44369586e-02,  2.88739172e-02,  1.82868142e-02,
  2.88739172e-02,  1.21270452e-01,  2.21366699e-02,
  6.73724735e-03,  4.13859480e-02,  1.63618864e-02,
  3.27237729e-02,  7.69971126e-03,  2.88739172e-02,
  1.18383061e-01,  1.82868142e-02,  4.81231954e-03,
  5.48604427e-02,  3.46487007e-01],
 [ 4.63718593e+01,  9.98418593e+01,  5.52763819e+00,
  7.98362548e+03,  2.65829146e+00,  6.89949749e+00,
  2.00391960e+03,  1.54291457e+02,  2.80502513e+01,
  7.88944724e-01,  2.11055276e-01,  3.33066907e-16,
  5.82914573e-01,  1.20603015e-01,  6.03015075e-02,
  1.00502513e-02,  5.02512563e-03,  7.53768844e-02,
  1.45728643e-01,  4.52261307e-02,  2.51256281e-02,
  5.02512563e-03,  1.00502513e-02,  3.51758794e-02,
  3.51758794e-02,  1.25628141e-01,  2.51256281e-02,
  1.00502513e-02,  3.51758794e-02,  2.01005025e-02,
  2.51256281e-02,  5.02512563e-03,  2.51256281e-02,
  1.15577889e-01,  2.51256281e-02,  5.02512563e-03,
  2.51256281e-02,  4.02010050e-01],
 [ 4.07491103e+01,  9.95422598e+01,  6.26690391e+00,
  5.30138568e+03,  2.73309609e+00,  7.12989324e+00,
  2.00380427e+03,  1.27149466e+02,  4.08665480e+01,
  2.08166817e-17,  1.00000000e+00, -6.66133815e-16,
  4.55516014e-01,  1.01423488e-01,  6.22775801e-02,
  3.91459075e-02,  1.95729537e-02,  1.61921708e-01,
  1.60142349e-01,  6.76156584e-02,  2.13523132e-02,
  1.24555160e-02,  1.95729537e-02,  2.13523132e-02,
  3.02491103e-02,  9.60854093e-02,  1.77935943e-02,
  5.33807829e-03,  3.02491103e-02,  1.77935943e-02,
  3.91459075e-02,  1.24555160e-02,  3.73665480e-02,
  1.26334520e-01,  2.13523132e-02,  1.77935943e-02,
  4.62633452e-02,  3.59430605e-01],
 [ 3.50762463e+01,  9.02900000e+01,  6.60312805e+00,
  3.42798675e+03,  2.71358749e+00,  7.06842620e+00,
  2.00380059e+03,  1.03577713e+02,  5.62355816e+01,
  2.08166817e-17,  7.62463343e-01,  2.37536657e-01,
  3.17693060e-01,  1.20234604e-01,  1.33919844e-01,
  1.14369501e-01,  1.85728250e-02,  1.01661779e-01,
  1.93548387e-01,  6.15835777e-02,  1.95503421e-02,
  9.77517107e-03,  2.63929619e-02,  2.44379277e-02,
  3.71456500e-02,  1.06549365e-01,  2.34604106e-02,
  3.91006843e-03,  4.49657869e-02,  2.05278592e-02,
  2.34604106e-02,  9.77517107e-03,  2.24828935e-02,
  1.22189638e-01,  2.05278592e-02,  1.46627566e-02,
  5.47409580e-02,  3.53861193e-01]])
```

```
In [40]: cc = pd.DataFrame(data=model.cluster_centers_, columns=[df_temp.columns])
```

```
In [41]: cc
```

Out[41]:

	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	MSRP	PRODUCTCODE	Large
0	29.889317	65.438037	6.618864	1884.194591	2.724735	7.132820	2003.815207	73.337825	63.256978	2.081668e-17
1	46.371859	99.841859	5.527638	7983.625477	2.658291	6.899497	2003.919598	154.291457	28.050251	7.889447e-01
2	40.749110	99.542260	6.266904	5301.385676	2.733096	7.129893	2003.804270	127.149466	40.866548	2.081668e-17
3	35.076246	90.290000	6.603128	3427.986755	2.713587	7.068426	2003.800587	103.577713	56.235582	2.081668e-17

4 rows × 38 columns

```
In [ ]:
```